

```
import os
import cv2
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import tensorflow as tf
from sklearn.model_selection import train_test_split
from tensorflow.keras import layers, models
from tensorflow.keras.applications import VGG16
from tensorflow.keras.layers import Dense, Flatten, Dropout
from tensorflow.keras.models import Model
from tensorflow.keras.optimizers import Adam
from sklearn.preprocessing import LabelEncoder
from tensorflow.keras.utils import to_categorical

pip install keras

Requirement already satisfied: keras in /usr/local/lib/python3.10/dist-packages (2.15.0)

from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive
```

## ✓ Load Dataset

```
# @title Load Dataset
def load_data(dataset_dir, img_size=(224, 224)):
    images=[]
    labels=[]

    classes=os.listdir(dataset_dir)
    label_encoder=LabelEncoder()
    for class_name in classes:
        class_dir=os.path.join(dataset_dir, class_name)

        if not os.path.isdir(class_dir):
            continue
        for img_file in os.listdir(class_dir):
            img_path=os.path.join(class_dir, img_file)

            img=cv2.imread(img_path)
            img=cv2.resize(img, img_size)
            img=cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

            images.append(img)
            labels.append(class_name)
    labels_encoded = label_encoder.fit_transform(labels)
    labels_one_hot = to_categorical(labels_encoded)

    return np.array(images), np.array(labels_one_hot)

dataset_dir='/content/drive/MyDrive/db/Indian Food Images'
images, labels=load_data(dataset_dir)

print("Number of images:", len(images))
print("Number of labels:", len(labels))

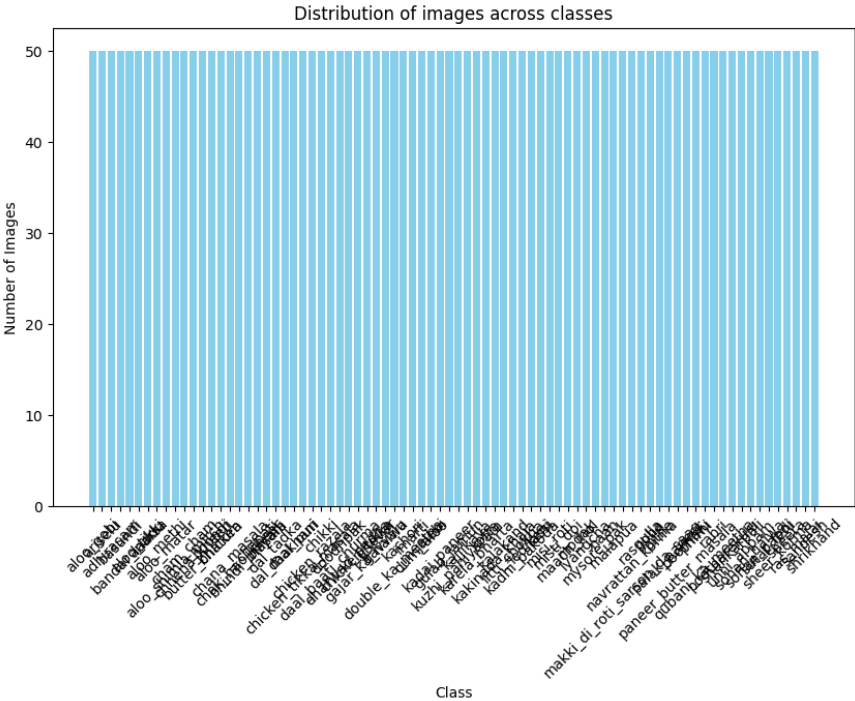
Number of images: 4000
Number of labels: 4000
```

## ✓ Data exploration

```
# @title Data exploration
classes=os.listdir(dataset_dir)

class_counts={}
for class_name in classes:
    class_dir=os.path.join(dataset_dir, class_name)
    num_images=len(os.listdir(class_dir))
    class_counts[class_name]=num_images
```

```
plt.figure(figsize=(10, 6))
plt.bar(class_counts.keys(), class_counts.values(), color='skyblue')
plt.xlabel('Class')
plt.ylabel('Number of Images')
plt.title('Distribution of images across classes')
plt.xticks(rotation=45)
plt.show()
```



```
class_df=pd.DataFrame(list(class_counts.items()), columns=['Class', 'Number of images'])
class_df
```

	Class	Number of images	
0	aloo_gobi	50	
1	ariselu	50	
2	adhirasam	50	
3	basundi	50	
4	bandar_laddu	50	
...	...	...	
75	sheer_korma	50	
76	sheera	50	
77	ras_malai	50	
78	sandesh	50	
79	shrikhand	50	

80 rows × 2 columns

Next steps:

Generate code with class\_df

☒ View recommended plots

```
total_imgs=sum(class_counts.values())
print("Total no of images:", total_imgs)
print("No of classes:", len(classes))
print("Avg imgs per class:", total_imgs/len(classes))

Total no of images: 4000
No of classes: 80
Avg imgs per class: 50.0
```

Preprocessing

```
# @title Preprocessing
preprocessed_images=[]
preprocessed_labels=[]

for img, label in zip(images, labels):
    img=img.astype(np.float32)/255.0

    preprocessed_images.append(img)
    preprocessed_labels.append(label)

preprocessed_images=np.array(preprocessed_images)
preprocessed_labels=np.array(preprocessed_labels)
print("Preprocessed images shape:", preprocessed_images.shape)
print("Preprocessed labels shape:", preprocessed_labels.shape)

Preprocessed images shape: (4000, 224, 224, 3)
Preprocessed labels shape: (4000, 80)
```

Transfer Learning

```
# @title Transfer Learning
def create_model(input_shape, num_classes):
    model=models.Sequential([
        layers.Conv2D(32, (3, 3), activation='relu', input_shape=input_shape),
        layers.MaxPooling2D((2, 2)),
        layers.Conv2D(64, (3, 3), activation='relu'),
        layers.MaxPooling2D((2, 2)),
        layers.Conv2D(128, (3, 3), activation='relu'),
        layers.MaxPooling2D((2, 2)),

        layers.Flatten(),

        layers.Dense(128, activation='relu'),
        layers.Dense(num_classes, activation='softmax')
    ])

    model.compile(optimizer='adam',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

    return model

input_shape=(224, 224, 3)
num_classes=80
model=create_model(input_shape, num_classes)

model.summary()

Model: "sequential"

```

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 222, 222, 32)	896
max_pooling2d (MaxPooling2D)	(None, 111, 111, 32)	0
conv2d_1 (Conv2D)	(None, 109, 109, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 54, 54, 64)	0
conv2d_2 (Conv2D)	(None, 52, 52, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 26, 26, 128)	0
flatten (Flatten)	(None, 86528)	0
dense (Dense)	(None, 128)	11075712
dense_1 (Dense)	(None, 80)	10320

```

Total params: 11179280 (42.65 MB)
Trainable params: 11179280 (42.65 MB)
Non-trainable params: 0 (0.00 Byte)
```

## Model Training

```
# @title Model Training
def build_transfer_learning_model(input_shape, num_classes):
    base_model=VGG16(weights='imagenet', include_top=False, input_shape=input_shape)

    for layer in base_model.layers:
        layers.trainable=False

    x=Flatten()(base_model.output)
    x=Dense(512, activation='relu')(x)
    x=Dropout(0.5)(x)
    output=Dense(num_classes, activation='softmax')(x)

    model=Model(inputs=base_model.input, outputs=output)

    model.compile(optimizer=Adam(),
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])
    return model

transfer_learning_model=build_transfer_learning_model(input_shape, num_classes)

transfer_learning_model.summary()
```

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels](https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels)  
58889256/58889256 [=====] - 4s 0us/step  
Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[None, 224, 224, 3]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
flatten_1 (Flatten)	(None, 25088)	0
dense_2 (Dense)	(None, 512)	12845568
dropout (Dropout)	(None, 512)	0
dense_3 (Dense)	(None, 80)	41040

=====  
Total params: 27601296 (105.29 MB)  
Trainable params: 27601296 (105.29 MB)  
Non-trainable params: 0 (0.00 Byte)

## Evaluation

```
# @title Evaluation
x_data, y_data=load_data(dataset_dir)

x_train, x_test, y_train, y_test=train_test_split(x_data, y_data, test_size=0.2, random_state=42)
print(x_train.shape)
print(y_train.shape)

epochs=10
batch_size=32

history=transfer_learning_model.fit(x_train, y_train,
                                    epochs=epochs,
                                    batch_size=batch_size,
                                    validation_data=(x_test, y_test))

(3200, 224, 224, 3)
(3200, 80)
Epoch 1/10
100/100 [=====] - 71s 435ms/step - loss: 5.7060 - accuracy: 0.0134 - val_loss: 4.3840 - val_accuracy: 0.0111
Epoch 2/10
100/100 [=====] - 44s 443ms/step - loss: 4.3820 - accuracy: 0.0109 - val_loss: 4.3856 - val_accuracy: 0.0089
Epoch 3/10
100/100 [=====] - 42s 415ms/step - loss: 4.3809 - accuracy: 0.0094 - val_loss: 4.3877 - val_accuracy: 0.0089
Epoch 4/10
100/100 [=====] - 44s 441ms/step - loss: 4.6207 - accuracy: 0.0131 - val_loss: 4.4374 - val_accuracy: 0.0111
Epoch 5/10
100/100 [=====] - 44s 440ms/step - loss: 4.4212 - accuracy: 0.0128 - val_loss: 4.3912 - val_accuracy: 0.0089
Epoch 6/10
100/100 [=====] - 42s 421ms/step - loss: 4.3752 - accuracy: 0.0181 - val_loss: 4.3934 - val_accuracy: 0.0111
Epoch 7/10
100/100 [=====] - 42s 422ms/step - loss: 4.3654 - accuracy: 0.0188 - val_loss: 4.3935 - val_accuracy: 0.0089
Epoch 8/10
100/100 [=====] - 42s 420ms/step - loss: 4.3560 - accuracy: 0.0213 - val_loss: 4.3962 - val_accuracy: 0.0089
Epoch 9/10
100/100 [=====] - 44s 438ms/step - loss: 4.3716 - accuracy: 0.0172 - val_loss: 4.3972 - val_accuracy: 0.0089
Epoch 10/10
100/100 [=====] - 42s 418ms/step - loss: 4.3522 - accuracy: 0.0206 - val_loss: 4.3992 - val_accuracy: 0.0089
```

Start coding or [generate](#) with AI.